

# Collective Communications

Professor Seokkoo Kang

Department of Civil & Environmental Engineering


[kangsk78@hanyang.ac.kr](mailto:kangsk78@hanyang.ac.kr)

# Collective Communication

- All process participate in the communication together. Must be called simultaneously.
- MPI\_Barrier
- MPI\_Bcast
- MPI\_Reduce
- MPI\_Allreduce
- MPI\_Gather
- MPI\_Scatter


# MPI\_Barrier

- MPI\_Barrier(MPI\_Comm communicator)  
Synchronize all processes



# MPI\_Bcast (Broadcast)

- Distribute the same information from the root process to all others
- `MPI_Bcast(void* data, int count, MPI_Datatype datatype, int root, MPI_Comm communicator)`



# MPI\_Bcast Example

```
MPI_Comm_size(MPI_COMM_WORLD,&size);
MPI_Comm_rank(MPI_COMM_WORLD,&rank);

if(rank==0){
    for(i=0;i<10;i++)
        buffer[i]=i*2;
}
MPI_Bcast(buffer, 10, MPI_INT, 0, MPI_COMM_WORLD);
if(rank!=0){
    for(i=0;i<10;i++)
        printf("Process %d: buffer[%d]=%d\n", rank, i, buffer[i]);
}
MPI_Finalize();
```

# Output

Process 1: buffer[0]=0

Process 1: buffer[1]=2

Process 1: buffer[2]=4

Process 1: buffer[3]=6

Process 1: buffer[4]=8

Process 1: buffer[5]=10

Process 1: buffer[6]=12

Process 1: buffer[7]=14

Process 1: buffer[8]=16

Process 1: buffer[9]=18

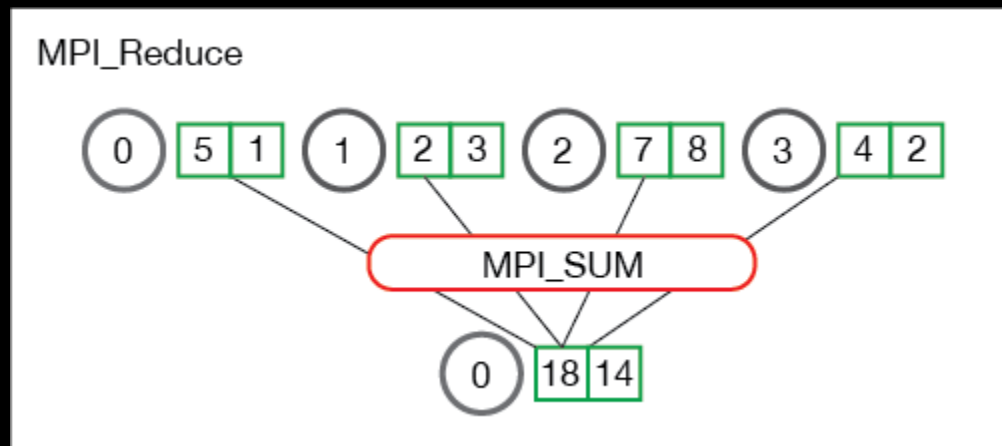
# MPI\_Reduce

**Reduces values on all processes to a single value**

```
int MPI_Reduce ( void *sendbuf, void *recvbuf, int count,  
                MPI_Datatype datatype, MPI_Op op, int root,  
                MPI_Comm comm    )
```

**MPI\_Op:** MPI\_SUM, MPI\_MAX, MPI\_MIN, MPI\_PROD

**Root:** the core # to which the result is sent.



# Sample MPI code

```
MPI_Comm_rank (MPI_COMM_WORLD, &rank);
MPI_Comm_size (MPI_COMM_WORLD, &size);
...
int my_sum=0;
int total_sum=0;
for (i=n_start; i < n_end; i++) my_sum = my_sum + i;
MPI_Reduce(&my_sum, &total_sum, 1, MPI_INT,
           MPI_SUM,0, MPI_COMM_WORLD);
if(rank==0) printf("#%d, Total sum=%d\n", total_sum);
if(rank==1) printf("#%d, Total sum=%d\n", total_sum);
...
```



# Output

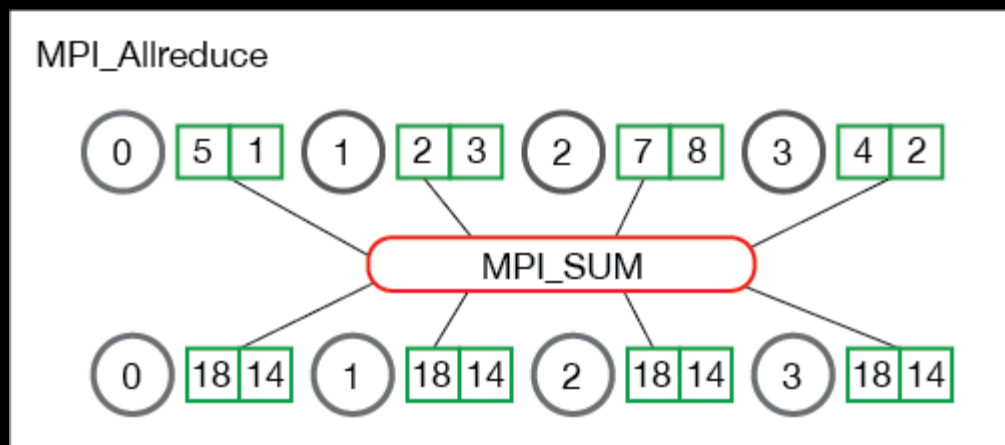


**#0, Total sum=5050**

**#1, Total sum=0**

# MPI\_Allreduce

## Reduce + Broadcast

```
int MPI_Allreduce ( void *sendbuf, void *recvbuf, int count,  
                   MPI_Datatype datatype, MPI_Op op, MPI_Comm comm )
```



# Sample MPI code

```
MPI_Comm_rank (MPI_COMM_WORLD, &rank);
MPI_Comm_size (MPI_COMM_WORLD, &size);
...
my_sum=0;
total_sum=0;
for (i=n_start; i < n_end; i++) my_sum = my_sum + i;
MPI_Allreduce(&my_sum, &total_sum, 1, MPI_INT,
              MPI_SUM, MPI_COMM_WORLD);
if(rank==0) printf("#%d, Total sum=%d\n", total_sum);
if(rank==1) printf("#%d, Total sum=%d\n", total_sum);
...
```


# Output

**#0, Total sum=5050**

**#1, Total sum=5050**

# MPI\_Scatter


- Splits the data from one process (=root) to all other processes (compare it with MPI\_Bcast)



**MPI\_Scatter** (void\* send\_data, int send\_count, MPI\_Datatype send\_datatype, void\* recv\_data, int recv\_count, MPI\_Datatype recv\_datatype, int root, MPI\_Comm communicator)

# MPI\_Gather

- Collects the data from each process and stores the data in process rank order in `recv_data` on the root.



**MPI\_Gather** (`void*` send\_data, `int` send\_count, `MPI_Datatype` send\_datatype, `void*` recv\_data, `int` recv\_count, `MPI_Datatype` recv\_datatype, `int` root, `MPI_Comm` communicator)

# MPI\_Scatter Example

senddata

1  
2  
3  
4  
5  
6  
7  
8  
9  
10

MPI\_SCATTER

recvdata

1  
2  
3  
4  
5  
6  
7  
8  
9  
10

```
totalrank=5; /*assume total process=5 */  
int senddata[10], recvdata[2];  
int sendsize=2, recvsize=2;  
if (rank==0) for(i=0; i<10; i++) senddata[i]=i+1;  
MPI_Scatter(senddata, sendsize, MPI_INT,  
            recvdata, recvsize, MPI_INT, 0,  
            MPI_COMM_WORLD);
```


# MPI\_Gather Sample Code

sum\_array

my\_sum

3  
7  
13  
15  
19

← MPI\_GATHER




```
int my_sum=0;
for(int i=0; i<recvsize; i++)
    my_sum = my_sum + recvddata[i];
int *sum_array = malloc( sizeof(int) *total_rank );
MPI_Gather ( &my_sum, 1, MPI_INT,
            &sum_array[0], 1, MPI_INT,
            0, MPI_COMM_WORLD);
```

**MPI\_Gather** ( void\* send\_data, int send\_count, MPI\_Datatype send\_datatype,  
void\* recv\_data, int recv\_count, MPI\_Datatype recv\_datatype,  
int root, MPI\_Comm communicator)



# Vairants

- MPI\_Scatterv: variable length
- MPI\_Gatherv
- MPI\_Allgather (MPI\_Gather+MPI\_Bcast)



# Assignment

- Write a code for calculating the sum from 1 to 10,000 and printing the result in two different ways.
  1. Use only MPI\_Scatter and MPI\_Gather.
  2. Use only MPI\_Reduce.
- Upload the ppt material after each presentation.